**Ajay Venkatesh**
Brooklyn, NY
+1 516 585 9013
av3855@nyu.edu

May 9, 2026

Hiring Manager
Cerebras Systems

Dear Hiring Manager,

My name is Ajay Venkatesh — finishing my M.S. in Computer Engineering at NYU, with several years of production software experience: a few at PwC in Bengaluru, a summer SDE internship at Amazon, and ongoing on-campus work at NYU's Novel AI Technologies. My focus has been on backend systems, cloud infrastructure, and the practical AI layer — **RAG pipelines, agentic workflows,** and infrastructure for **ML inference at scale**.

The Cerebras MTS posting maps cleanly to where I've been pushing my work. At Amazon I provisioned resilient **Infrastructure-as-Code** on **AWS CDK** across auto-scaling **ECS** clusters with multi-region CI/CD, and built a high-throughput, event-driven triage system on **ECS Fargate, SQS, and SNS** that materially reduced response latency under variable load. The shape of that work — containerized, autoscaling, cost-aware orchestration of inference-style workloads — is the same family of problem the posting describes for high-performance, low-latency inference services.

On the AI side, I deployed a **Retrieval-Augmented Generation** pipeline with an **agentic workflow** at Amazon, and implemented a **Model Context Protocol (MCP)** server that materially reduced Athena query time and compute costs through tool-augmented retrieval. I've also built RAG systems on **AWS Bedrock** with FAISS indexing and shipped **Python** pipelines for data preprocessing, model integration, and inference execution.

Stack-wise: Python is my daily driver for both backend services and ML pipelines; I'm fluent in **Docker** and **Kubernetes** for orchestration; comfortable in **C++** from graduate coursework and systems projects; and I'm at home debugging distributed systems through logs, metrics, and distributed traces. My graduate coursework — **Computer Architecture, Machine Learning, Deep Learning, Big Data, System Design** — gave me the systems-and-ML intuition that Cerebras's stack rewards.

What pulls me toward Cerebras is the conviction behind the WSE — that inference throughput and latency are bottlenecked by general-purpose hardware, and the answer is a stack rebuilt from silicon up. I'd like to work on that stack at the inference-service layer, alongside the people building it.

Thanks for considering my application — I'd welcome the chance to discuss further.

Sincerely,
Ajay Venkatesh