

Ajay Venkatesh

Brooklyn, NY

+1 516 585 9013

av3855@nyu.edu

May 13, 2026

Hiring Manager

Tavern Research

Dear Hiring Manager,

My name is Ajay Venkatesh — finishing my M.S. in Computer Engineering at NYU, based here in Brooklyn, with several years of production software experience: a few years at PwC in Bengaluru, a summer SDE internship at Amazon, and ongoing on-campus work at NYU’s Novel AI Technologies. My focus has been on backend systems, the practical AI layer — LLM-powered features, RAG pipelines, and agentic workflows — and shipping things end-to-end in small, fast-moving teams.

The Engineer I posting reads like a description of how I already work. **Python** is my daily driver for both backend and ML pipelines; **JavaScript / TypeScript** is where I spend my time on the product surface; I’m fluent in **Docker** and **Kubernetes** for containerization and deployment; and I’ve shipped APIs and database-backed systems across multiple stacks. At Amazon I deployed a **Retrieval-Augmented Generation (RAG)** pipeline with an **agentic workflow** that autonomously triages incidents by reasoning over operational logs, and implemented a **Model Context Protocol (MCP)** server for tool-augmented retrieval. At Campus Mesh I built **CampusIQ**, an LLM-powered AI assistant with iterative **prompt engineering** and **RAG retrieval** over structured data. Deploying LLMs and AI-assisted systems in production isn’t a side project — it’s been the shape of my last year of work.

On using AI to accelerate development specifically: I use **Cursor** and **GitHub Copilot** daily, but I verify trade-offs and understand what’s underneath before shipping production code. Fundamentals plus AI, not either alone — which I think is what the Engineer I description is getting at when it pairs “ship clean code” with “use LLMs to expand your skill set.”

On backend systems and reliability: at PwC I spent three years on a distributed backend on **Microsoft Azure** integrating with **Microsoft CRM**, with event-driven processing, fault-tolerant retries over SQL Server, and parallelized batch ingestion that cut execution from hours to minutes. At Amazon I provisioned **Infrastructure-as-Code** on **AWS CDK** (Lambda, API Gateway, DynamoDB, IAM, CloudWatch) with multi-stage CI/CD and rollback controls. I’ve debugged production incidents through logs, metrics, and traces — the actual unglamorous work of keeping a system reliable.

What pulls me toward Tavern is the combination: a small team in Brooklyn, modern stack, AI-native workflows, and a product space where the work actually matters — civic engagement isn’t a synthetic problem. I’d like to work somewhere I can keep building at speed, contribute across the stack, and learn from senior engineers without the friction of a big-company process.

I’d welcome the chance to discuss further. Thanks for considering my application.

Sincerely,

Ajay Venkatesh