

Ajay Venkatesh

Brooklyn, NY

+1 516 585 9013

av3855@nyu.edu

May 13, 2026

Hiring Manager

TEAM GmbH

Dear Hiring Manager,

My name is Ajay Venkatesh — finishing my M.S. in Computer Engineering at NYU, with several years of production software experience: a few years at PwC in Bengaluru, a summer SDE internship at Amazon, and ongoing on-campus work at NYU's Novel AI Technologies. My focus has been on **backend systems**, **SQL-driven data workflows**, and the practical **AI layer** — tooling that actually helps engineers ship faster on top of stable, production-grade systems.

The ProStore® posting maps onto a few of my strongest threads. Most of my professional work has lived in **Java**, **Python**, and **SQL** — and the SQL side specifically has been the load-bearing layer of every project I've shipped. At PwC I spent three years on a distributed backend on **Microsoft Azure** integrating with **Microsoft CRM**, with event-driven processing and fault-tolerant retries over **SQL Server** stores — writing complex queries, stored logic, and parallelized batch ingestion that cut execution time from hours to minutes. The shift from SQL Server to **Oracle SQL / PL/SQL** is a natural next step rather than a different discipline, and I'd come in expecting the dialect-specific learning curve.

On warehouse and operations-style systems specifically: at Amazon I architected a secure console for fulfillment center teams to query image and video evidence across FCs — understanding the messy, customer-specific shape of logistical workflows and translating them into a system operators could trust. I provisioned the supporting infrastructure with **AWS CDK** (API Gateway, Lambda, DynamoDB, IAM, CloudWatch) and shipped backend services in **Java** (Coral, Smithy, Dagger, CBOR) with unit/integration testing via **Mockito**.

On AI in the development process: I use **Cursor** and **GitHub Copilot** daily, and I've shipped AI-assisted tooling — a **Retrieval-Augmented Generation** pipeline with an **agentic workflow** at Amazon for operational log reasoning, plus a **Model Context Protocol (MCP)** server for tool-augmented retrieval. The role's framing of "AI to support the development process" is exactly where I keep pushing my own workflow.

What pulls me toward TEAM GmbH is the seriousness of the product space — warehouse management is one of those domains where stability and trust beat novelty every time. I'd like to work somewhere the customer's logistics problems are real, the data lives in Oracle for a reason, and the engineering team values careful, documented solutions over flashy ones.

I'd welcome the chance to discuss further. Thanks for considering my application.

Sincerely,

Ajay Venkatesh