

Education

New York University | *Master of Science in Computer Engineering (GPA – 3.9/4)*

Sep 2024 – May 2026

- *Coursework:* System Design, Machine Learning, Deep Learning, Big Data, Advanced Java, Data Structures

Skills

Programming Languages: Python, Java, C++, TypeScript, JavaScript (ES6), SQL, Ruby, Scala, Swift, Golang

Web Technologies: REST API, GraphQL API, Microservices, React, Node.js, Express.js, Sequelize

Database: RDBMS (MySQL, PostgreSQL), NoSQL (Dynamo DB, FireBase, MongoDB), Dataverse

Cloud: AWS (AWS CDK, EC2, S3, Dynamo DB, API Gateway, Lambda, Cognito, IAM, Q, SageMaker), Azure (OpenAI services, Function Apps, Security Data lakes, Storage blobs)

Tools: Git, Visual Studio, Cursor, Jupyter, Postman, Docker, Power Automate, Power BI, QuickSight, MapReduce, HDFS, Spark, Hadoop, Pig, Hive, SproutCore

AI/ML: PyTorch, Hugging Face, LLM Fine-tuning, RAG (LangChain, FAISS), NLP, SVM, Random Forest, MLflow, PySpark

Experience

CampusMesh

Jan 2026 – Present

Software Engineer | *Real-time microservices, recommendation engine, and LLM-based AI retrieval pipelines*

- Developed scalable **Node.js microservices** powering real-time academic workflows across study groups, messaging, calls, notifications, and profile management, backed by **MySQL (Sequelize)** and **DynamoDB**.
- Deployed **serverless architecture** on AWS using **Serverless Framework** with **Lambda, API Gateway, S3, and CloudFront** across **4 environments** (local, dev, staging, production); enforced security with **KMS-managed keys, SSM Parameter Store, and least-privilege IAM**.
- Worked on **LLM-based retrieval pipelines (RAG)** for CampusIQ (AI Assistant), leveraging **embedding-based semantic search** and structured data indexing to improve contextual accuracy by **30%**.

Amazon

May 2025 – Aug 2025

Software Development Engineer Intern | *Serverless AWS infrastructure, React 18 frontend, and Java backend systems*

- Architected a secure UI console for Amazon fulfillment center teams to view image/video evidence across FCs; presented the **system design** for approval and established Role Based Access (**4+ roles**) using **OAuth, Midway SSO, and HELIS**.
- Provisioned **serverless Infrastructure-as-Code** using **AWS CDK** with API Gateway, Lambda, DynamoDB, S3, and IAM, reducing provisioning time by **70%** and enabling **multi-stage, multi-region CI/CD deployments**.
- Built a **React 18** frontend with optimized API integration, advanced filtering, client-side caching, reducing page load time by **40%**.
- Engineered backend services in **Java** using **Coral, Smithy, Dagger**, and **CBOR serialization**; applied **low-level design and design patterns** reduced backend execution time by **35%** and ensure reliability by unit and integration testing using **Mockito**.

New York University (Novel AI Technologies, Inc.)

Aug 2024 – Present

Software Developer (On-Campus) | *NSF-funded AI health platform, ETL pipelines, and full-stack insurance analytics*

- Architected an **NSF-funded (\$100K)** AI-powered health app for **iOS** and **Android**, processing real-time wearable data; deployed cloud-hosted AI models for health metrix analysis and process food images.
- Engineered ETL pipelines transforming **NoSQL** into **PostgreSQL** using **concurrency control and locking mechanisms**; processed **250K+** records and powered **AWS QuickSight** dashboards for real-time health analytics.
- Developed an insurance analytics platform (**React, Node.js, REST APIs**); deployed (**Docker** containers on **EC2**, implemented Stripe billing, RBAC, and row-level security, generating **\$50K+** revenue.
- Streamlined underwriting via broker portal integrated with insurer platform; implemented **document and media processing** and **WebSocket-based real-time communication**; secured **6 insurance clients**.

PricewaterhouseCoopers (PwC)

Aug 2021 – Aug 2024

Senior Software Engineer | *Cloud-native backend, NLP/OCR automation, and data engineering on Azure*

- Developed a distributed backend on **Microsoft Azure** for a health platform, integrating with **Microsoft CRM** and implementing event-driven processing and fault tolerance.
- Led code reviews and mentored engineers on **software design principles**, reducing production bugs by **25%**.
- Deployed an **NLP based OCR neural model** on Azure to automate structured data extraction, reducing data processing time by **75%**.
- Optimized data ingestion with **parallel processing** for **100K-row** from Excel; reduced execution time from **5 hours to 1 hour**.
- Built **Power BI dashboards** processing large datasets from multiple databases, improving reporting efficiency by **40%**.

Projects

Flight Delay Prediction (Big Data ML) | *PySpark, HDFS, MapReduce, Random Forest, Gradient Boosted Tree*

- Built a scalable **PySpark ML pipeline** on **HDFS** to process and engineer features from over **7M flight records** using SQL-based aggregation and anomaly detection.
- Performed extensive EDA to identify delay patterns; trained and optimized **Random Forest** and **Gradient Boosted Tree** models achieving **85%+** accuracy.
- Converted model outputs into actionable insights, enabling data-driven scheduling and operational planning for delay reduction.

Certifications

AWS Cloud Practitioner	Azure Data Fundamentals	Microsoft Solution Architect Expert
Power BI Data Analyst Associate	Microsoft Developer Associate	Stanford Machine Learning (Coursera)