

Education

New York University, New York, NY

Sep 2024 – May 2026

Master of Science in Computer Engineering (GPA: 3.9/4.0)

Coursework: System Design, Machine Learning, Deep Learning, Big Data, Advanced Java, Data Structures

Experience

- **Amazon** May 2025 – Aug 2025
Software Development Engineer Intern Seattle, US
 - Engineered high-throughput backend services in **Java** (Coral, Smithy, Dagger, CBOR) across a **distributed architecture**; applied **concurrency-aware design patterns** with unit/integration testing via **Mockito**.
 - Provisioned **serverless infrastructure** using **AWS CDK** with API Gateway, Lambda, **DynamoDB**, S3, and IAM; supported **multi-stage CI/CD** deployments and production debugging for live services.
 - Built a performant **React** internal console with optimized API integration and caching; triaged defects with testers in a production-adjacent environment.
- **CampusMesh** Jan 2026 – Present
Software Engineer New York, US
 - Developed scalable **Node.js microservices** in a **distributed system** backed by **MySQL (Sequelize)** and **DynamoDB**, supporting real-time messaging, notifications, and profile workflows.
 - Deployed services on **Linux/AWS** (Lambda, API Gateway, S3) with **Terraform** and **Serverless Framework** across local through production environments.
- **New York University (Novel AI Technologies, Inc.)** Aug 2024 – Present
Software Developer (On-Campus) New York, US
 - Engineered **PostgreSQL** data pipelines with **concurrency control and locking**, transforming **NoSQL** sources into relational stores for analytics and reporting workloads.
 - Deployed backend services via **Docker** on **Linux/EC2**; implemented **RBAC** and row-level security for production data access.
- **PricewaterhouseCoopers (PwC)** Aug 2021 – Aug 2024
Senior Software Engineer Bengaluru, India
 - Developed a **distributed backend** on **Microsoft Azure** with asynchronous message processing and fault-tolerant retries over **SQL Server** in a live production environment.
 - Diagnosed complex production incidents across multi-service flows; led code reviews and optimized bulk data ingestion with parallel execution on large datasets.

Projects

- **Advanced Expense Tracker (Java Systems Application)** | *Java, JDBC, Maven, Multithreading, OOP*
 - Architected a modular **Java application** with **OOP design patterns**, **JDBC**, and **multithreading/synchronization** for thread-safe **concurrency control** over shared database state.
- **InviLite (Distributed Data Pipeline)** | *Python, AWS, Lambda, DynamoDB, Kinesis, S3*
 - Built a cloud-native ingestion pipeline (Kinesis, Lambda, S3, Glue, **DynamoDB**) for structured and unstructured data with governed storage primitives and failure triage workflows.

Skills

- **Programming Languages:** Java, C, C++, Go, Python, SQL
- **Systems & Backend:** Distributed Systems, Concurrency Control, REST API, Microservices, Linux
- **Database:** PostgreSQL, MySQL, DynamoDB, SQL Server, MongoDB, JDBC
- **Cloud & DevOps:** AWS, Kubernetes, Terraform, Docker, CI/CD, Azure
- **Tools:** Git, Mockito, JUnit, Postman, Debugging, Cursor, GitHub Copilot

Certifications

AWS Cloud Practitioner

Azure Data Fundamentals