

Education

New York University, New York, NY

Aug 2024 – May 2026

Master of Science in Computer Engineering (GPA: 3.9/4.0)

Coursework: System Design, Machine Learning, Deep Learning, Big Data, Data Structures

Experience

- **CampusMesh** Jan 2026 – Present
Software Engineer
New York, US
 - Developed scalable **Python/Node.js microservices** in a **large-scale distributed architecture** for real-time messaging, notifications, and profile workflows, backed by **MySQL** and **DynamoDB**.
 - Built **embedding-based semantic search** and data indexing for CampusIQ (AI Assistant), supporting data ingestion, retrieval, and evaluation-style introspection over structured campus datasets.
 - Provisioned **cloud infrastructure** on **AWS** (Lambda, API Gateway, S3, CloudFront) with **CI/CD** across local through production environments; enforced security with **KMS**, **SSM**, and least-privilege **IAM**.
- **Amazon** May 2025 – Aug 2025
Software Development Engineer Intern
Seattle, US
 - Provisioned **serverless Infrastructure-as-Code** using **AWS CDK** with API Gateway, Lambda, DynamoDB, S3, and IAM; reduced provisioning time **70%** with **multi-stage, multi-region CI/CD**.
 - Engineered backend services in a **distributed microservice architecture**; tuned API and backend performance, cutting execution time by **35%** with unit/integration testing via **Mockito**.
 - Built internal tooling and a **React** console for operational debugging and production-adjacent issue triage across fulfillment center workflows.
- **New York University (Novel AI Technologies, Inc.)** Aug 2024 – Present
Software Developer
New York, US
 - Engineered **Python ETL/data processing pipelines** transforming **NoSQL** into **PostgreSQL** with **concurrency control and locking**; processed **250K+** records for analytics and reporting workloads.
 - Deployed **ML inference services** and backend workloads via **Docker** on **Linux/EC2** for health metric analysis and wearable data processing.
- **PricewaterhouseCoopers (PwC)** Aug 2021 – Aug 2024
Senior Software Engineer
Bengaluru, India
 - Developed a **distributed backend** on **Microsoft Azure** with asynchronous message processing and fault-tolerant retries over **SQL Server** stores in a production environment.
 - Optimized bulk data ingestion with **parallel processing** on **100K-row** datasets, cutting transformation runtime from **5 hours to 1 hour**; diagnosed multi-service production incidents via logs and metrics.

Projects

- **InviLite (Large-Scale Data Processing Pipeline)** | *Python, AWS, Kinesis, Lambda, S3, Glue, DynamoDB, FAISS*
 - Architected a cloud-native **Python data platform** (Kinesis, Lambda, S3, Glue, DynamoDB) for governed ingestion, storage, and mining of structured and unstructured enterprise data at scale.
 - Built an **ML evaluation pipeline** (SVM + FAISS + Bedrock LLM) with measurable classification improvements and failure-analysis workflows for production triage.
- **Advanced Expense Tracker (Python Systems Application)** | *Python, Multithreading, SQLite3, Pandas*
 - Implemented concurrent transaction handling with **Python threading** and locks for thread-safe writes; applied **OOP design patterns** and performance-conscious data access over **SQLite3**.

Skills

- **Programming Languages:** Python, C++, Go, Java, SQL, TypeScript
- **Systems & Backend:** Distributed Systems, Microservices, Multithreading, REST API, Linux
- **Data & ML:** ETL Pipelines, PostgreSQL, DynamoDB, MongoDB, PyTorch, RAG (FAISS), NLP
- **Cloud & DevOps:** AWS (CDK, Lambda, EC2, S3, API Gateway, DynamoDB, IAM), Docker, Kubernetes, CI/CD
- **Tools & Platforms:** Git, Cursor, GitHub Copilot, Postman, Mockito, Jupyter

Certifications

[AWS Cloud Practitioner](#)

[Azure Data Fundamentals](#)