

Ajay Venkatesh

Brooklyn, NY

+1 516 585 9013

av3855@nyu.edu

June 11, 2026

Hiring Manager

Goldman Sachs

New York, New York

Dear Hiring Manager,

I'm Ajay Venkatesh, completing an M.S. in Computer Engineering at NYU, with production software experience at PwC in Bengaluru, a summer SDE internship at Amazon, and ongoing on-campus engineering work at NYU's Novel AI Technologies. I'm applying for the Data Engineer role on your Lakehouse and AI data platform, where I'd help build reliable batch and streaming pipelines, curated datasets, and operational controls that support analytics and AI use cases.

At NYU, I engineered **Python** and **ANSI SQL** ETL workflows into **PostgreSQL** with thoughtful **schema design**, concurrency controls, and validation checks across **250K+** records. I troubleshoot pipeline breaks, optimized SQL for reporting consumers, and partnered with operations stakeholders to deliver documented, usable data products backed by **AWS QuickSight** dashboards.

In projects, I built a distributed **Apache Spark/PySpark** pipeline over **7M+** records with **Spark SQL**, **Parquet**-friendly processing, and **data quality** checks from raw ingestion through curated feature tables. I also architected a monitored ingestion workflow on **AWS (Kinesis, Lambda, S3, Glue, DynamoDB, CloudWatch)** that strengthened day-to-day operational support. At Amazon, I shipped event-driven services on **SQS/SNS** with tested **Java** backends and **CI/CD** releases; at PwC, I modernized batch feeds on **Azure/SQL Server** and shaped reporting datasets for business consumers.

My stack includes **Python, Java, SQL, Apache Spark, Databricks, Kafka, JSON, Parquet, Kubernetes, Docker, AWS, GCP, and Git**. **Snowflake, Apache Iceberg, Avro, and Sybase IQ** are platforms I have not owned end-to-end in production yet, but my **Spark**, lakehouse-style pipeline, and cloud data delivery background give me a strong base to ramp quickly on Goldman Sachs' internal tooling.

What draws me to Goldman Sachs is the chance to work on production data platforms where reliability, reconciliation discipline, and clear standards for testing and monitoring directly affect how teams trust analytics and AI outputs. I care about collaborating with engineers and data consumers, communicating risks early, and improving shared platform tooling where it strengthens delivery.

I'd welcome the chance to discuss further. Thanks for considering my application.

Sincerely,

Ajay Venkatesh