

## Skills

**Programming Languages:** Python, Java, SQL (ANSI SQL), JSON

**Data Engineering:** Apache Spark, PySpark, Databricks, ETL/ELT, Data Pipelines, Batch & Streaming, Schema Design, Schema Evolution, Temporal Data Modelling, Data Compatibility, Clustering, Data Quality

**Databases & Formats:** PostgreSQL, SQL Server, MySQL, DynamoDB, MongoDB, Parquet

**Cloud & Platforms:** AWS (S3, Glue, Lambda, Kinesis, Redshift, QuickSight, CDK, CloudWatch), GCP (BigQuery), Azure (Data Lake, DevOps)

**Streaming & Orchestration:** Kafka, Event-Driven Ingestion, CI/CD, Kubernetes, Docker, Git

**Tools:** Jupyter, Postman, Power BI, Visual Studio, Cursor

## Experience

**New York University (Novel AI Technologies, Inc.)**

**Aug 2024 – Present**

*Data Engineer*

*New York, US*

- Engineered **Python/SQL ETL pipelines** into **PostgreSQL** with **schema design**, concurrency control, and locking; processed **250K+** records and validated completeness and consistency for analytics and reporting consumers.
- Built curated datasets and **AWS QuickSight** dashboards for operations partners; troubleshoot pipeline breaks, optimized **ANSI SQL** queries, and documented data products for cross-team use.
- Deployed containerized data services on **Docker/AWS EC2**; refactored ingestion paths to improve reliability, maintainability, and operational support across multi-tenant workloads.

**Amazon**

**May 2025 – Aug 2025**

*Software Development Engineer Intern*

*Seattle, US*

- Supported production-ready backend **Java** services with **unit/integration tests**; provisioned **AWS CDK** infrastructure (**Lambda, DynamoDB, S3, CloudWatch**) with **multi-stage CI/CD** across regions.
- Built an event-driven platform on **AWS ECS Fargate, SQS, and SNS** with structured logging and monitoring; diagnosed cross-service defects and incorporated review feedback before release.

**PricewaterhouseCoopers (PwC)**

**Aug 2021 – Aug 2024**

*Data Engineer*

*Bengaluru, India*

- Developed batch ingestion and distributed processing on **Microsoft Azure** over **SQL Server**; modernized repetitive feeds with **parallel processing**, cutting runtime from **5 hours to 1 hour**.
- Built **Power BI** reporting datasets joined across relational sources; partnered with business consumers to shape usable outputs and improve delivery consistency.

**CampusMesh**

**Jan 2026 – Present**

*Data Engineer*

*New York, US*

- Maintained **Node.js** services and **MySQL/DynamoDB** stores across **4 environments**; enforced data-access controls with **KMS, SSM**, and least-privilege **IAM** on **AWS**.

## Projects

**Flight Delay Prediction (Lakehouse-Style Analytics Pipeline) | *PySpark, Apache Spark, Spark SQL, HDFS, Parquet, ANSI SQL***

- Built a distributed **Apache Spark/PySpark** pipeline ingesting **7M+** flight records; applied **SQL** aggregations, feature engineering, and **data quality** checks for analytics and ML outputs.
- Trained **Spark MLlib** models on cleaned datasets; documented pipeline stages from raw ingestion through curated feature tables.

**InviLite (Streaming Data & Monitoring Pipeline) | *Python, AWS Kinesis, Lambda, S3, Glue, DynamoDB, CloudWatch***

- Architected ingestion workflows for structured and unstructured enterprise data into governed stores; added **CloudWatch** monitoring and automated triage, cutting incident resolution from **1–2 hours to minutes**.

## Education

**New York University** | *Master of Science in Computer Engineering (GPA – 3.9/4)*

**Sep 2024 – May 2026**

- *Coursework:* Big Data, System Design, Machine Learning, Deep Learning, Data Structures

## Certifications

**AWS Cloud Practitioner**

**Azure Data Fundamentals**

**Stanford Machine Learning (Coursera)**